

HPE CRAY PROGRAMMING ENVIRONMENT

Isambard Hackathon

March 23, 2021

John M. Levesque

John.levesque@hpe.com

Technical Advisor for Cray Programming Environment

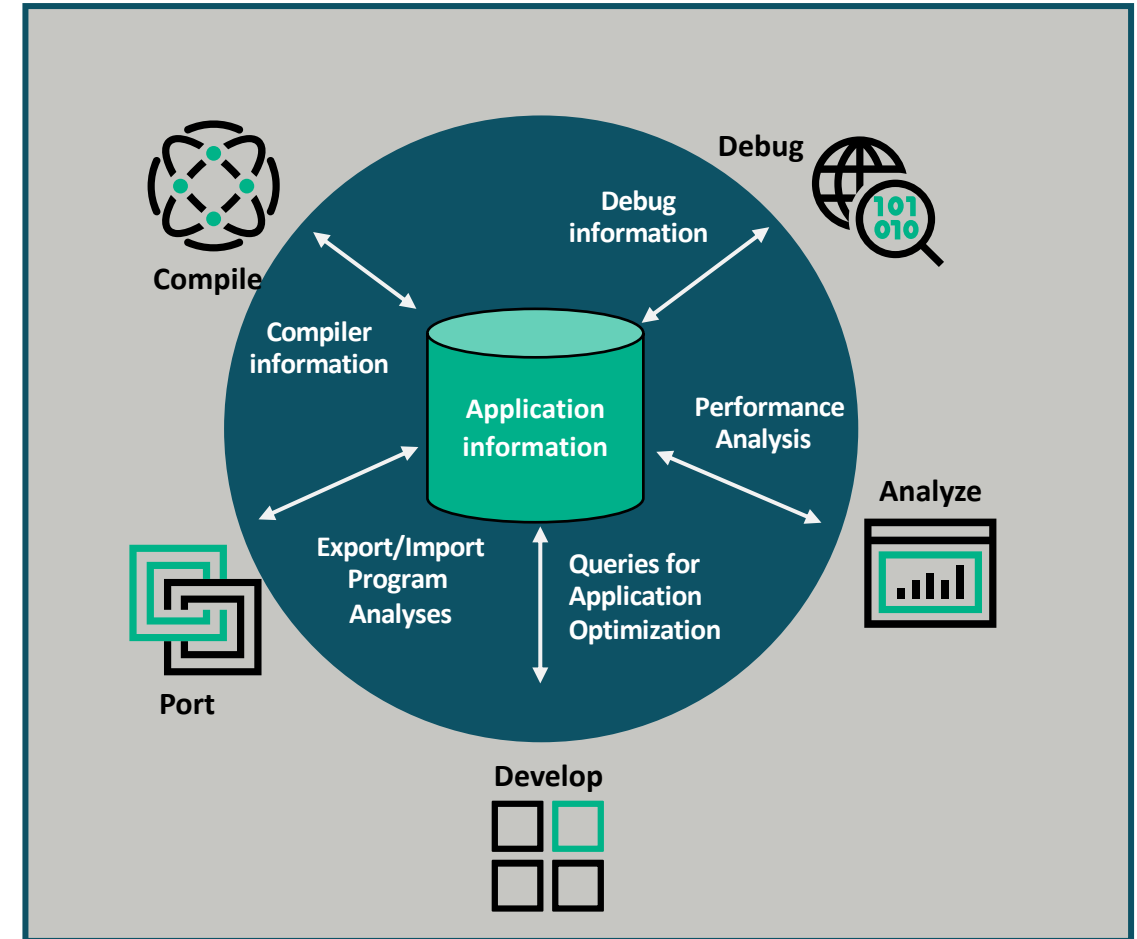
Distinguished Technologist

CTO Office

CRAY PROGRAMMING ENVIRONMENT




Why Cray PE?

- Full, production quality developer software stack that targets performance and programmability at scale
- Minimized time to productive use of system
- Focuses on real HPC applications, not just benchmarks
- Issues addressed in a timely fashion
- Traditional HPC capability coupled with advancing technology to embrace new problems
- Engagement with customers to understand goals and challenges




CRAY PE TECHNOLOGY APPLIED TO ARM SVE

- Cray technology **designed for real scientific applications**, not just for benchmarks
- **Modules simplify build environment**
- **Fortran, C, and C++ compilers**
 - Most complete vectorization capabilities that have evolved from custom Cray vectorizing processors, including unique outer loop vectorization technology
 - SVE instruction set has many features familiar to Cray's compiler such as wider vector widths, predication, gather/scatter, etc.
- Scalable debug tools like **gdb4hpc** and **CCDB** assist with porting to new processors or platforms
- **Scientific Libraries** extract maximum performance using standard interfaces for ease of use
- **Cray Reveal** combines compiler knowledge of the application with profiling information
 - Unique tool in industry that significantly reduces effort associated with adding OpenMP
 - Scoping tool to help users port and optimize applications
- **Cray Performance tools** deliver wealth of profiling capability including whole-program view of performance with simple-to-use interfaces



COMPILING ENVIRONMENT ADVANTAGES

We designed our compilers for real-life applications, not just benchmarks



Performance and programmability

- Our compiler exploits the scalar and vector hardware capabilities of the systems
- Compiler optimization feedback for application tuning

Fully integrated heterogeneous optimization capability

- Providing consistency across all HPE and HPE Cray systems, supporting:
 - X86-64 (both Intel & AMD) processors
 - Arm-based processors
 - GPUs

Integration with tools for performance optimization

- Integration with parallelization assistant through Program Library technology for adding OpenMP to your applications
- Integration with performance analysis tools for additional optimization insight

Focus on application portability and investment protection

Focus on compliance with latest standards and language support:

- Languages: Fortran with coarrays, C/C++ and UPC
- Programming models: OpenMP and OpenACC



CRAY COMPILER FOR BAYMAX

- Supports native compilation (no cross-compiling)
- Supports OpenMP, Cray Reveal
- Offers loopmark (compiler feedback)

- cce-sve
 - Fortran and C/C++ compiler generates ARM SVE code
 - C/C++ compiler based on Cray classic compiler (EDG front-end)
 - Not as strong C++ support

- cce
 - Released in Sept to provide stronger compiler for C++ code
 - Fortran and C/C++ compiler generates ARM Neon vector code
 - C/C++ compiler based on new Cray clang compiler (LLVM)
 - ARM SVE code generation coming when it's available in LLVM (targeting spring of 2021)
 - Will move to SVE code generation for Fortran and C/C++

- Usage guidance
 - Choose compiler based on dominant/most important code (C++ strength vs Fortran SVE strength)
 - Cannot mix use between two compiling environments
 - Libsci targets Fujitsu A64FX processor and can be used with either cce-sve or cce compilers



MODULES FOR RUNNING WITH SVE AND PERFTOOLS

Currently Loaded Modulefiles:

- | | |
|----------------------------|---|
| 1) cpe-cray | 5) craype-network-infiniband |
| 2) cce-sve/10.0.2(default) | 6) cray-libsci/20.10.1.2(default) |
| 3) craype/2.7.4(default) | 7) slurm/slurm/no-version |
| 4) craype-arm-nsp1 | 8) perftools-base/21.02.0(default) |
| 9) perftools-lite | 10) cray-mvapich2_nogpu_noregcache_sve/2.3.5(default) |

MODULES FOR RUNNING WITH SVE WITHOUT PERFTOOLS

Currently Loaded Modulefiles:

- | | |
|-----------------------------|--|
| 1) cpe-cray | 5) craype-network-infiniband |
| 2) cce-sve/10.0.2 (default) | 6) cray-libsci/20.10.1.2 (default) |
| 3) craype/2.7.4 (default) | 7) slurm/slurm/no-version |
| 4) craype-arm-nsp1 | 8) cray-mvapich2_nogpu_sve/2.3.5 (default) |

LESLIE3D

- Three Dimension Computational Fluid Dynamics Code from Georgia Tech



SIMPLE PROFILE

Perftools-lite places profiles into standard out. - all exclusive times

Table 1: Profile by Function

Samp%	Samp	Imb.	Imb.	Group
		Samp	Samp%	Function=[MAX10]
				PE=HIDE
100.0%	3,586.7	--	--	Total

68.7%	2,465.6	--	--	USER

19.2%	687.9	40.1	5.7%	fluxk_.LOOP@li.38
17.0%	611.0	19.0	3.1%	fluxi_.LOOP@li.30
13.7%	493.1	11.9	2.4%	fluxj_.LOOP@li.38
4.6%	165.8	13.2	7.7%	extrapk_.LOOP@li.158
3.0%	106.8	10.2	9.1%	update_.LOOP@li.18
2.8%	99.5	15.5	13.9%	extrapj_.LOOP@li.159
=====				
28.6%	1,024.1	--	--	MPI

11.2%	401.4	74.6	16.3%	MPI_SEND
9.7%	348.3	131.7	28.5%	MPI_ALLREDUCE
5.3%	191.1	221.9	55.7%	MPI_REDUCE
=====				
2.6%	94.8	--	--	ETC

2.2%	78.7	61.3	45.4%	__ALLOCATE
=====				

Sampling run 1/100
sec

Column

- 1 Percent
- 2 # samples
- 3 #imblanced samples
- 4 Percent of imbalance
- 5 Code segment

LINE LEVEL PROFILE

Perftools-lite places profiles into standard out

Table 2: Profile by Group, Function, and Line

Samp%	Samp	Imb.	Imb.	Group
		Samp	Samp%	Function=[MAX10]
				Source
				Line
				PE=HIDE
100.0%	3,586.7	--	--	Total

68.7%	2,465.6	--	--	USER

19.2%	687.9	--	--	fluxk_.LOOP@li.38
3				levesque/leslie3d_mpi_orig/src/fluxk.f

4	1.4%	51.5	12.5	20.3% line.41
4	13.8%	495.5	31.5	6.2% line.75
4	2.4%	85.6	30.4	27.2% line.84
=====				
17.0%	611.0	--	--	fluxi_.LOOP@li.30
3				levesque/leslie3d_mpi_orig/src/fluxi.f

4	14.0%	502.3	20.7	4.1% line.65
4	1.0%	37.0	12.0	25.4% line.72

Sampling run 1/100
sec

Column

- 1 Percent
- 2 # samples
- 3 #imblanced samples
- 4 Percent of imbalance
- 5 Code segment
- 6 Line level

ALWAYS USE -HLIST=A TO GET COMPILER LISTING

```
38. + F-----<      DO J = 1, JCMAX
39. + F F-----<      DO K = 0, KCMAX
40.   F F V-----<      DO I = 1, IND
41.   F F V              QS(I) = UAV(I,J,K) * SKX(I,J,K) +
42.   F F V              >      VAV(I,J,K) * SKY(I,J,K) +
43.   F F V              >      WAV(I,J,K) * SKZ(I,J,K)
44.   F F V
45.   F F V              IF ( NSCHEME .EQ. 2 ) THEN
46.   F F V                L = K + 1 - KADD
47.   F F V                QSP = U(I,J,L) * SKX(I,J,K) +
48.   F F V              >      V(I,J,L) * SKY(I,J,K) +
49.   F F V              >      W(I,J,L) * SKZ(I,J,K)
50.   F F V                QSPK = (QSP - QS(I)) * DBLE(1 - 2 * KADD)
51.   F F V                IF (QSPK .GT. 0.0D+00) QS(I) = 0.5D+00 * (QS(I) + QSP)
52.   F F V              ENDIF
53.   F F V
54.   F F V              FSK(I,K,1) = QAV(I,J,K,1) * QS(I)
55.   F F V              FSK(I,K,2) = QAV(I,J,K,2) * QS(I) +
56.   F F V              >      PAV(I,J,K) * SKX(I,J,K)
57.   F F V              FSK(I,K,3) = QAV(I,J,K,3) * QS(I) +
58.   F F V              >      PAV(I,J,K) * SKY(I,J,K)
59.   F F V              FSK(I,K,4) = QAV(I,J,K,4) * QS(I) +
60.   F F V              >      PAV(I,J,K) * SKZ(I,J,K)
61.   F F V              FSK(I,K,5) = (QAV(I,J,K,5) + PAV(I,J,K)) *
62.   F F V              >      QS(I)
63.   F F V
64.   F F V              IF (ISGSK .EQ. 1) THEN
65.   F F V                FSK(I,K,7) = QAV(I,J,K,7) * QS(I)
66.   F F V              ENDIF
67.   F F V
68.   F F V              IF ( ICHM .GT. 0 ) THEN
69.   F F V D-----<      DO L = 8, 7 + NSPECI
70.   F F V D              FSK(I,K,L) = QAV(I,J,K,L) * QS(I)
71.   F F V D----->      ENDDO
72.   F F V              ENDIF
73.   F F V----->      ENDDO
74.   F F
75.   F F V I-----<      IF ( VISCOUS ) CALL VISCK ( 1, IND, J, K, FSK )
76.   F F----->      ENDDO
```

NUMA INFORMATION

Perftools-lite places profiles into standard out

Table 3: Memory Bandwidth by Numanode

Memory Traffic GBytes	Read Memory Traffic GBytes	Write Memory Traffic GBytes	Thread Time	Memory Traffic GBytes / Sec	Memory Traffic / Nominal Peak	Numanode Node Id=[max3,min3] PE=HIDE
736.29	596.75	139.54	36.205149	20.34	7.9%	numanode.0
736.29	596.75	139.54	36.205149	20.34	7.9%	nid.19
731.44	592.71	138.73	36.156605	20.23	7.9%	nid.18
718.71	582.49	136.22	36.149776	19.88	7.8%	nid.17
703.97	570.05	133.92	36.156420	19.47	7.6%	nid.22
703.83	569.91	133.91	36.196569	19.44	7.6%	nid.20
689.34	558.02	131.32	36.157255	19.07	7.4%	nid.23



SAMPLING WITH CALL TREE

To get call-tree pat_report –Oct <experiment directory>

Table 1: Function Calltree View

Samp%	Samp	Calltree
		PE=HIDE
100.0%	3,586.7	Total

79.9%	2,867.5	les3d_

23.8%	855.0	fluxk_
3 23.8%	854.9	fluxk_.REGION@li.33 (ovhd)

4 19.2%	688.6	fluxk_.REGION@li.33
5 19.2%	688.0	fluxk_.LOOP@li.38
4 4.6%	166.0	extrapk_
5		extrapk_.REGION@li.156 (ovhd)
6 4.6%	165.8	extrapk_.REGION@li.156
7		extrapk_.LOOP@li.158
=====		
18.8%	674.5	fluxi_
3 18.8%	674.4	fluxi_.REGION@li.25 (ovhd)

4 17.0%	611.5	fluxi_.REGION@li.25
5 17.0%	611.1	fluxi_.LOOP@li.30
4 1.7%	62.7	extrapi_
5		extrapi_.REGION@li.137 (ovhd)
6 1.7%	62.4	extrapi_.REGION@li.137
7 1.2%	42.9	extrapi_.LOOP@li.141
=====		

NO MPI STATISTICS WITH PERFTOOLS-LITE, SO LETS TRY PERFTOOLS

Module swap perftools-lite perftools

Compile and link

Pat_build -u -g mpi <executable>

Srun run executable+pat



SIMPLE PROFILE

Pat_report <experiment file>

Table 1: Profile by Function Group and Function

Time%	Time	Imb.	Imb.	Calls	Group
		Time	Time%		Function
					PE=HIDE
100.0%	35.764334	--	--	36,574.1	Total

71.4%	25.550017	--	--	5,074.6	USER

19.3%	6.895061	0.140679	2.1%	200.0	fluxk_.LOOP@li.38
17.1%	6.112560	0.137001	2.3%	200.0	fluxi_.LOOP@li.30
13.7%	4.902728	0.074435	1.6%	200.0	fluxj_.LOOP@li.38
4.6%	1.639311	0.087245	5.2%	200.0	extrapk_.LOOP@li.158
3.0%	1.062270	0.598870	37.4%	1,010.0	mpicx_
2.8%	1.004442	0.025128	2.5%	200.0	extrapj_.LOOP@li.159
1.6%	0.578741	0.013319	2.3%	21.0	tmstep_
1.5%	0.545074	0.011006	2.1%	100.0	update_
1.4%	0.508350	0.037538	7.1%	202.0	parallel_
1.4%	0.501630	0.011941	2.4%	100.0	update_.LOOP@li.18
1.3%	0.457291	0.039064	8.2%	476.1	ghost_
1.1%	0.409303	0.017826	4.3%	200.0	extrapi_.LOOP@li.141
=====					

LOAD IMBALANCE WITH MPI

Pat_report <experiment file>

Table 2: Profile of maximum function times

Time%	Time	Imb.	Imb.	Function
		Time	Time%	PE=[max,min]

100.0%	7.035740	0.140679	2.1%	fluxk_.LOOP@li.38

100.0%	7.035740	--	--	pe.8
95.9%	6.745043	--	--	pe.21
=====				
88.8%	6.249561	0.137001	2.3%	fluxi_.LOOP@li.30

88.8%	6.249561	--	--	pe.4
85.2%	5.991107	--	--	pe.25
=====				
70.7%	4.977162	0.074435	1.6%	fluxj_.LOOP@li.38

70.7%	4.977162	--	--	pe.16
68.4%	4.812507	--	--	pe.3
=====				
64.6%	4.542432	0.697049	15.9%	MPI_SEND

64.6%	4.542432	--	--	pe.0
48.5%	3.413080	--	--	pe.20
=====				

LOAD IMBALANCE AND MESSAGE SIZES WITH MPI

Pat_report <experiment file>

Table 3: Load Balance with MPI Message Stats

Time%	Time	MPI Msg Count	MPI Msg Bytes	Avg MPI Msg Size	Group PE=[mmm]
100.0%	35.764426	9,968.9	1,765,373,643.4	177,088.87	Total
71.4%	25.550109	0.0	0.0	--	USER
78.2%	27.961406	0.0	0.0	--	pe.0
71.2%	25.479947	0.0	0.0	--	pe.9
68.5%	24.497551	0.0	0.0	--	pe.25
28.5%	10.194769	9,968.9	1,765,373,643.4	177,088.87	MPI
31.5%	11.247938	8,526.0	1,530,361,072.0	179,493.44	pe.25
28.5%	10.186792	8,526.0	1,540,057,072.0	180,630.67	pe.3
21.8%	7.783913	8,526.0	1,540,057,072.0	180,630.67	pe.0

MESSAGE SIZES – MIN, MAX, AVE WITH MPI

Pat_report <experiment file>

```
-----  
MPI Msg Bytes%                54.2%  
MPI Msg Bytes                956,672,000.0  
MPI Msg Count                10,100.0 msgs  
16<= MsgSz <256 Count        0.0 msgs  
256<= MsgSz <4KiB Count      0.0 msgs  
4KiB<= MsgSz <64KiB Count    6,060.0 msgs  
64KiB<= MsgSz <1MiB Count    4,040.0 msgs  
16MiB<= MsgSz Count          0.0 msgs  
=====
```

MPI_SEND / mpicx_ / parallel_ / les3d_ / pe.18

```
-----  
MPI Msg Bytes%                53.6%  
MPI Msg Bytes                946,976,000.0  
MPI Msg Count                10,100.0 msgs  
16<= MsgSz <256 Count        0.0 msgs  
256<= MsgSz <4KiB Count      0.0 msgs  
4KiB<= MsgSz <64KiB Count    6,060.0 msgs  
64KiB<= MsgSz <1MiB Count    4,040.0 msgs  
16MiB<= MsgSz Count          0.0 msgs  
=====
```

MPI_SEND / mpicx_ / parallel_ / les3d_ / pe.27

```
-----  
MPI Msg Bytes%                35.3%  
MPI Msg Bytes                623,776,000.0  
MPI Msg Count                8,080.0 msgs  
16<= MsgSz <256 Count        0.0 msgs  
256<= MsgSz <4KiB Count      0.0 msgs  
4KiB<= MsgSz <64KiB Count    6,060.0 msgs  
64KiB<= MsgSz <1MiB Count    2,020.0 msgs  
=====
```

MESSAGE SIZES – MIN, MAX, AVE WITH MPI

Pat_report –Oct <experiment file>

Table 1: Function Calltree View

Time%	Time	Calls	Calltree
			PE=HIDE
100.0%	35.764334	--	Total

99.5%	35.572095	1.0	les3d_

23.9%	8.547970	--	fluxk_

19.3%	6.902694	200.0	fluxk_.REGION@li.33
19.3%	6.895061	200.0	fluxk_.LOOP@li.38
4.6%	1.642442	--	extrapk_
4.6%	1.640225	200.0	extrapk_.REGION@li.156
4.6%	1.639311	200.0	extrapk_.LOOP@li.158
=====			
18.9%	6.751193	--	fluxi_

17.1%	6.117161	200.0	fluxi_.REGION@li.25
17.1%	6.112560	200.0	fluxi_.LOOP@li.30
1.8%	0.631296	--	extrapi_
1.8%	0.628479	200.0	extrapi_.REGION@li.137
1.1%	0.409303	200.0	extrapi_.LOOP@li.141
=====			
18.5%	6.629379	202.0	parallel_

15.8%	5.663738	1,010.0	mpicx_

10.8%	3.845383	9,522.9	MPI_SEND
3.0%	1.062270	1,010.0	mpicx_(exclusive)
1.9%	0.670306	9,522.9	MPI_WAIT
=====			

LETS LOOK AT LOOP INSTRUMENTATION

- Module swap perftools perftools-lite-loops
- DO EVER DO THIS WHEN RUNNING FOR REAL _ MAKE CODE RUN SLOWER



MESSAGE SIZES – MIN, MAX, AVE WITH MPI

Perftools-lite-loops place profile into standard out

Table 1: Inclusive and Exclusive Time in Loops (from compiler-inserted loop instrumentation)

Loop Incl Time%	Loop Incl Time	Time (Loop Adj.)	Loop Exec Trips	Loop Trips Avg	Loop Trips Min	Loop Trips Max	Function=/.LOOP[.] PE=HIDE

96.0%	62.483595	0.001207	1	100.0	100	100	les3d_.LOOP.3.li.218
89.8%	58.428053	12.470257	100	2.0	2	2	les3d_.LOOP.4.li.274
18.9%	12.303912	0.000742	200	96.0	96	96	fluxk_.LOOP.1.li.38
17.7%	11.551030	0.000195	200	27.4	27	28	fluxi_.LOOP.1.li.30
17.7%	11.550835	7.489365	5,486	96.0	96	96	fluxi_.LOOP.2.li.31
15.3%	9.965937	0.000450	200	27.4	27	28	fluxj_.LOOP.1.li.38
14.0%	9.103169	8.076826	19,200	28.4	28	29	fluxk_.LOOP.2.li.39
12.0%	7.804776	7.170560	5,486	97.0	97	97	fluxj_.LOOP.2.li.39
10.8%	7.058184	0.007562	202	5.0	5	5	parallel_.LOOP.1.li.16
8.6%	5.625883	5.625883	526,629	97.0	97	97	visci_.LOOP.1.li.809
8.5%	5.515949	5.515949	545,829	96.0	96	96	visck_.LOOP.1.li.367
8.5%	5.515182	0.000325	200	28.4	28	29	extrapk_.LOOP.1.li.160
8.5%	5.514857	4.410280	5,686	99.0	99	99	extrapk_.LOOP.2.li.162
8.3%	5.375106	0.000202	200	31.1	30	32	extrapj_.LOOP.1.li.161
8.3%	5.374903	4.204186	6,229	97.0	97	97	extrapj_.LOOP.2.li.162
7.2%	4.688454	4.688454	532,114	96.0	96	96	viscj_.LOOP.1.li.367
5.7%	3.739388	0.000092	200	31.1	30	32	extrapi_.LOOP.1.li.141
5.7%	3.739296	3.739296	6,229	99.0	99	99	extrapi_.LOOP.2.li.142
5.3%	3.470722	1.246789	526,629	96.0	96	96	fluxi_.LOOP.4.li.69
4.9%	3.200002	0.023447	19,200	27.4	27	28	fluxk_.LOOP.4.li.79
4.9%	3.176555	0.457949	526,629	96.0	96	96	fluxk_.LOOP.5.li.81
4.2%	2.718606	2.718606	50,556,343	5.0	5	5	fluxk_.LOOP.6.li.83
3.9%	2.519885	2.519885	557,229	100.0	100	100	exk4_.LOOP.1.li.280
3.4%	2.223934	2.223934	50,556,343	5.0	5	5	fluxi_.LOOP.5.li.71
3.4%	2.205782	2.205782	597,943	100.0	100	100	exj4_.LOOP.1.li.280

Sampling run 1/100
sec

Column

- 1 Inclusive loop percent
- 2 Inclusive loop time
- 3 ?????
- 4 Time loop is executed
- 5 Avg trip count
- 6 Min trip count
- 7 Max trip count
- 8 Location of loop

MESSAGE SIZES – MIN, MAX, AVE WITH MPI (PAT_REPORT -OCT

Table 1: Calltree with Loop Inclusive Time

Incl	Incl	Loop Exec	Loop	Calltree
Time%	Time		Trips	PE=HIDE
			Avg	
100.0%	65.09	--	--	Total

84.4%	54.97	--	--	les3d_

96.0%	62.48	1	100.0	les3d_.LOOP.3.li.218

89.8%	58.43	100	2.0	les3d_.LOOP.4.li.274

17.6%	11.45	--	--	fluxk_

18.9%	12.30	200	96.0	fluxk_.LOOP.1.li.38

14.0%	9.10	19,200	28.4	fluxk_.LOOP.2.li.39

1.6%	1.03	545,829	96.0	fluxk_.LOOP.3.li.40
=====				
4.9%	3.20	19,200	27.4	fluxk_.LOOP.4.li.79

4.9%	3.18	526,629	96.0	fluxk_.LOOP.5.li.81

4.2%	2.72	50,556,343	5.0	fluxk_.LOOP.6.li.83
=====				
5.3%	3.42	--	--	exk4_

3.9%	2.52	557,229	100.0	exk4_.LOOP.1.li.280
=====				

Sampling run 1/100
sec

Column

- 1 Inclusive percent
- 2 Inclusive time
- 3 Number of times
element
encountered
- 4 Ave trip count
- 5 Program element

HOW WOULD WE USE PERFTOOLS TO PARALLELIZE THIS APPLICATION?



WHEN USING PERFTOOL-LITE TOOLS WITH C++

- Your initial profile will probably be sparse
 - `pat_report -P <statistics_file>`
- Perftools-lite-loops turns off OpenMP, if the application calls openMP API calls, they will be ignored

USING PERFTOOLS-LITE (OR -LOOPS OR -HBM)

- module load perftools-lite or perftools-lite-loops or perftools-lite-hbm
 - Module perftools-base should already be loaded
- Build application
- Run application
- Statistics report comes out within standard out
 - Also generates a directory of profile data to be examined with different options

PERFTOOLS-LITE PROFILE – RUN ON 8 NODES–1 MPI TASK/NODE

Table 1: Profile by Function

Samp%	Samp	Imb.	Imb.	Group
		Samp	Samp%	Function=[MAX10]
				PE=HIDE
100.0%	4,061.6	--	--	Total
95.0%	3,859.9	--	--	USER
21.4%	869.6	8.4	1.1%	fluxj_
20.7%	842.8	7.2	1.0%	fluxi_
19.9%	808.2	5.8	0.8%	fluxk_
7.9%	318.9	6.1	2.2%	extrapk_
7.5%	303.2	7.8	2.8%	extrapi_
6.8%	274.2	3.8	1.5%	update_
6.1%	249.4	4.6	2.1%	extrapj_
1.0%	40.5	5.5	13.7%	mpicx_
4.1%	165.2	--	--	MPI
2.2%	91.1	5.9	6.9%	MPI_REDUCE
1.1%	44.9	31.1	46.8%	MPI_SEND

Exclusive time
Sampling is in
100th of a second

Imbalance

Only showing items that
take up more than 1% of
time – you can override
with -T



PERFTOOLS-LITE PROFILE – RUN ON 8 NODES–8 MPI TASKS

Table 2: Profile by Group, Function, and Line

Samp%	Samp	Imb.	Imb.	Group
		Samp	Samp%	Function=[MAX10]
				Source
				Line
				PE=HIDE
100.0%	4,061.6	--	--	Total

95.0%	3,859.9	--	--	USER

21.4%	869.6	--	--	fluxj_
3				Leslie_CUG/leslie3d_mpi/src/fluxj.f

4	1.6%	63.6	8.4	13.3% line.31
4	14.0%	567.5	15.5	3.0% line.67
4	2.7%	109.6	12.4	11.6% line.76
=====				
20.7%	842.8	--	--	fluxi_
3				Leslie_CUG/leslie3d_mpi/src/fluxi.f

4	1.9%	77.1	18.9	22.5% line.24
4	14.2%	578.8	21.2	4.0% line.56
4	2.1%	85.5	10.5	12.5% line.63
=====				
19.9%	808.2	--	--	fluxk_
3				Leslie_CUG/leslie3d_mpi/src/fluxk.f

4	1.7%	69.0	13.0	18.1% line.31
4	12.2%	497.5	21.5	4.7% line.65
4	3.2%	129.6	18.4	14.2% line.74

Profile on line level within a routine – lets look at this routine



PERFTOOLS-LITE PROFILE – RUN ON 8 NODES–8 MPI TASKS

```
28. + F-----<      DO K = 1, KCMAX
29. + F F-----<      DO J = 0, JCMAX
30.   F F V-----<      DO I = 1, IND
31.     F F V          QS(I) = UAV(I,J,K) * SJX(I,J,K) +
32.     F F V          >          VAV(I,J,K) * SJY(I,J,K) +
33.     F F V          >          WAV(I,J,K) * SJZ(I,J,K)
34.     F F V
35.     F F V          IF (NSCHEME .EQ. 2) THEN
36.     F F V          L = J + 1 - JADD
37.     F F V
38.     F F V          QSP = U(I,L,K) * SJX(I,J,K) +
39.     F F V          >          V(I,L,K) * SJY(I,J,K) +
40.     F F V          >          W(I,L,K) * SJZ(I,J,K)
41.     F F V          QSPJ = (QSP - QS(I)) * DBLE(1 - 2 * JADD)
42.     F F V          IF (QSPJ .GT. 0.0D+00) QS(I) = 0.5D+00 * (QS(I) + QSP)
43.     F F V
44.     F F V          ENDIF
45.     F F V
46.     F F V          FSJ(I,J,1) = QAV(I,J,K,1) * QS(I)
47.     F F V          FSJ(I,J,2) = QAV(I,J,K,2) * QS(I) +
48.     F F V          >          PAV(I,J,K) * SJX(I,J,K)
49.     F F V          FSJ(I,J,3) = QAV(I,J,K,3) * QS(I) +
50.     F F V          >          PAV(I,J,K) * SJY(I,J,K)
51.     F F V          FSJ(I,J,4) = QAV(I,J,K,4) * QS(I) +
52.     F F V          >          PAV(I,J,K) * SJZ(I,J,K)
53.     F F V          FSJ(I,J,5) = (QAV(I,J,K,5) + PAV(I,J,K)) *
54.     F F V          >          QS(I)
55.     F F V
56.     F F V          IF (ISGSK .EQ. 1) THEN
57.     F F V          FSJ(I,J,7) = QAV(I,J,K,7) * QS(I)
58.     F F V          ENDIF
59.     F F V
60.     F F V          IF ( ICHEM .GT. 0 ) THEN
61.     F F V D-----<      DO L = 8, 7 + NSPECI
62.     F F V D          FSJ(I,J,L) = QAV(I,J,K,L) * QS(I)
63.     F F V D----->      ENDDO
64.     F F V          ENDIF
65.     F F V----->      ENDDO
66.     F F
67. + F F V I-----<>      IF ( VISCOS ) CALL VISCJ (1, IND, J, K, FSJ)
68.     F F----->      ENDDO
```

+ indicates further information

You get this listing by using
-hlist-a

Line 67 from -hlist=a

I indicates the call was
inlined , V loop Vectorized,
F Flattened



PERFTOOLS-LITE PROFILE – RUN ON 8 NODES–8 MPI TASKS

ftn-6315 ftn: VECTOR FLUXJ, File = fluxj.f, Line = 28

A loop starting at line 28 was not vectorized because the target array (qs) would require rank expansion.

ftn-3182 ftn: IPA FLUXJ, File = fluxj.f, Line = 28, Column = 7

Loop has been flattened.

ftn-6315 ftn: VECTOR FLUXJ, File = fluxj.f, Line = 29

A loop starting at line 29 was not vectorized because the target array (qs) would require rank expansion.

ftn-3182 ftn: IPA FLUXJ, File = fluxj.f, Line = 29, Column = 10

Loop has been flattened.

ftn-6204 ftn: VECTOR FLUXJ, File = fluxj.f, Line = 30

A loop starting at line 30 was vectorized.

ftn-6002 ftn: SCALAR FLUXJ, File = fluxj.f, Line = 61

A loop starting at line 61 was eliminated by optimization.

ftn-6383 ftn: VECTOR FLUXJ, File = fluxj.f, Line = 67

A loop starting at line 67 requires an estimated 25 vector registers at line 67; 2 of these have been preemptively forced to memory.

ftn-6204 ftn: VECTOR FLUXJ, File = fluxj.f, Line = 67

A loop starting at line 67 was vectorized.



PERFTOOLS-LITE PROFILE – RUN ON 8 NODES–8 MPI TASKS

Table 1: Function Calltree View

Samp%	Samp	Calltree
		PE=HIDE
100.0%	4,061.6	Total

85.1%	3,456.6	les3d_

27.8%	1,127.1	fluxk_

3 19.9%	808.2	fluxk_(exclusive)
3 7.9%	318.9	extrapk_
=====		
27.6%	1,119.0	fluxj_

3 21.4%	869.6	fluxj_(exclusive)
3 6.1%	249.4	extrapj_
=====		
20.7%	842.8	fluxi_
2.8%	115.4	parallel_
3 2.1%	87.0	mpicx_

4 1.1%	44.9	MPI_SEND
4 1.0%	40.5	mpicx_(exclusive)
=====		
2.3%	92.0	flowio_
3 2.2%	91.1	MPI_REDUCE
1.1%	45.9	tmstep_
=====		
7.5%	303.2	extrapi_
6.8%	274.2	update_

Pat_report –Oct <statistics directory>

Level in the Call Tree, everything 3 and higher is called from this || (2)



PERFTOOLS-LITE-LOOPS – RUN ON 8 NODES–8 MPI TASKS

Table 1: Inclusive and Exclusive Time in Loops (from -hprofile_generate)

Loop	Loop Incl	Loop Hit	Loop	Loop	Loop	Function=/.LOOP[.]
Incl	Time		Trips	Trips	Trips	PE=HIDE
Time%			Avg	Min	Max	

96.9%	60.109652	1	100.0	100	100	les3d_.LOOP.3.1i.216
96.0%	59.527700	100	2.0	2	2	les3d_.LOOP.4.1i.272
23.3%	14.461074	200	96.0	96	96	fluxi_.LOOP.1.1i.21
23.3%	14.460550	19,200	96.0	96	96	fluxi_.LOOP.2.1i.22
20.5%	12.724507	200	96.0	96	96	fluxk_.LOOP.1.1i.28
20.1%	12.486995	200	96.0	96	96	fluxj_.LOOP.1.1i.28
15.1%	9.370239	19,200	97.0	97	97	fluxj_.LOOP.2.1i.29
14.0%	8.713997	19,200	97.0	97	97	fluxk_.LOOP.2.1i.29
10.1%	6.259900	1,843,200	97.0	97	97	visci_.LOOP.1.1i.782
8.8%	5.445447	1,862,400	96.0	96	96	viscj_.LOOP.1.1i.347
8.0%	4.990874	1,862,400	96.0	96	96	visck_.LOOP.1.1i.348
7.9%	4.905076	200	97.0	97	97	extrapk_.LOOP.1.1i.141
7.9%	4.904410	19,400	99.0	99	99	extrapk_.LOOP.2.1i.143
7.6%	4.704280	1,843,200	96.0	96	96	fluxi_.LOOP.4.1i.60
6.6%	4.121449	200	99.0	99	99	extrapj_.LOOP.1.1i.141
6.6%	4.121028	19,800	97.0	97	97	extrapj_.LOOP.2.1i.142
6.5%	4.009192	19,200	96.0	96	96	fluxk_.LOOP.4.1i.69
6.4%	3.995474	1,843,200	96.0	96	96	fluxk_.LOOP.5.1i.71
5.5%	3.382283	200	99.0	99	99	extrapi_.LOOP.1.1i.123
5.5%	3.381946	19,800	99.0	99	99	extrapi_.LOOP.2.1i.124
5.0%	3.115988	19,200	96.0	96	96	fluxj_.LOOP.4.1i.71
5.0%	3.102463	1,843,200	96.0	96	96	fluxj_.LOOP.5.1i.73

This Table shows most important loops, the columns are percent of time, inclusive time, number of times the loop was executed, Avg, Min, Max iteration counts and location within the source



HOW DO I KNOW WHAT THE IMPORTANT LOOPS ARE?

- Pat_report -O calltree < directory produced by perftools-lite-loops run>
- Produces call tree with DO loops included

PERFTOOLS-LITE-LOOPS – RUN ON 8 NODES–8 MPI TASKS

Table 1: Function Calltree View

Time%	Time	Calls	Calltree
			PE=HIDE
100.0%	60.437329	--	Total

100.0%	60.437286	2.0	les3d_

96.8%	58.504830	--	les3d_.LOOP.3.li.216
3 95.8%	57.923177	--	les3d_.LOOP.4.li.272

30.9%	18.681501	200.0	fluxi_

12.1%	7.338945	200.0	fluxi_(exclusive)
11.4%	6.863554	--	fluxi_.LOOP.1.li.21
			fluxi_.LOOP.2.li.22
11.4%	6.863554	1,843,200.0	visci_
7.4%	4.479002	200.0	extrapi_

3.8%	2.320974	200.0	extrapi_(exclusive)
3.6%	2.158027	--	extrapi_.LOOP.1.li.123
			extrapi_.LOOP.2.li.124
3.6%	2.158027	1,960,200.0	exi4_
=====			
28.3%	17.098408	200.0	fluxk_

11.3%	6.809059	200.0	fluxk_(exclusive)
9.4%	5.654640	--	fluxk_.LOOP.1.li.28
			fluxk_.LOOP.2.li.29
9.4%	5.654640	1,862,400.0	visck_
7.7%	4.634709	200.0	extrapk_

4.9%	2.959637	--	extrapk_.LOOP.1.li.141
			extrapk_.LOOP.2.li.143
4.9%	2.937827	1,900,800.0	exk4_
2.8%	1.675072	200.0	extrapk_(exclusive)
=====			
26.6%	16.077506	200.0	fluxj_

10.2%	6.136293	--	fluxj_.LOOP.1.li.28
			fluxj_.LOOP.2.li.29
10.2%	6.136293	1,862,400.0	viscj_
10.1%	6.090073	200.0	fluxj_(exclusive)
6.4%	3.851140	200.0	extrapj_

USING REVEAL

- Need program library and perftools-lite-loops output
 - `ftn -hlist=a -hpl=leslie3d.pl`
 - `reveal leslie3d.pl < perftools-lite-loops data directory>`

DO NOT HAVE any PERFTOOLS-LITE or
PERFTOOLS modules loaded when you build
the program library

PERFTOOLS-LITE-LOOPS – RUN ON 8 NODES–1 MPI TASK/NODE

Table 1: Function Calltree View

Time%	Time	Calls	Calltree
			PE=HIDE
100.0%	60.437329	--	Total

100.0%	60.437286	2.0	les3d_

96.8%	58.504830	--	les3d_.LOOP.3.li.216
395.8%	57.923177	--	les3d_.LOOP.4.li.272

4330.9%	18.681501	200.0	fluxi_

511112.1%	7.338945	200.0	fluxi_(exclusive)
511111.4%	6.863554	--	fluxi_.LOOP.1.li.21
61111			fluxi_.LOOP.2.li.22
711111.4%	6.863554	1,843,200.0	visci_
511117.4%	4.479002	200.0	extrapi_

611113.8%	2.320974	200.0	extrapi_(exclusive)
611113.6%	2.158027	--	extrapi_.LOOP.1.li.123
71111			extrapi_.LOOP.2.li.124
811113.6%	2.158027	1,960,200.0	exi4_
=====			
411128.3%	17.098408	200.0	fluxk_

511111.3%	6.809059	200.0	fluxk_(exclusive)
511119.4%	5.654640	--	fluxk_.LOOP.1.li.28
61111			fluxk_.LOOP.2.li.29
711119.4%	5.654640	1,862,400.0	visck_
511117.7%	4.634709	200.0	extrapk_

611114.9%	2.959637	--	extrapk_.LOOP.1.li.141
71111			extrapk_.LOOP.2.li.143
811114.9%	2.937827	1,900,800.0	exk4_
611112.8%	1.675072	200.0	extrapk_(exclusive)
=====			
411126.6%	16.077506	200.0	fluxj_

5111110.2%	6.136293	--	fluxj_.LOOP.1.li.28
61111			fluxj_.LOOP.2.li.29
7111110.2%	6.136293	1,862,400.0	viscj_
5111110.1%	6.090073	200.0	fluxj_(exclusive)
511116.4%	3.851140	200.0	extrapj_

The screenshot shows the Reveal application window titled 'leslie3d.pl'. The interface includes a menu bar (File, Edit, View, Help) and a 'Navigation' pane on the left. The 'Navigation' pane displays a list of loops under the 'Loop Performance' section, each with a time value, a name, and a star icon. The loops listed are:

- 60.1097 LES3D@216
- 59.5277 LES3D@272
- 14.4611 FLUXI@21
- 14.4605 FLUXI@22
- 12.7245 FLUXK@28
- 12.4870 FLUXJ@28
- 9.3702 FLUXJ@29
- 8.7140 FLUXK@29
- 6.2599 FLUXI@782
- 5.4454 FLUXJ@347
- 4.9909 FLUXK@348
- 4.9051 FLUXK@141
- 4.9044 FLUXK@143
- 4.7043 FLUXI@60
- 4.1214 FLUXJ@141
- 4.1210 FLUXJ@142
- 4.0092 FLUXK@69
- 3.9955 FLUXK@71
- 3.3823 FLUXI@123
- 3.3819 FLUXI@124

Red arrows point from the 'Navigation' pane to the corresponding entries in the 'Table 1: Function Calltree View' on the left. The 'Source' pane on the right is empty, and the 'Info' pane at the bottom shows the status: 'leslie3d.pl loaded. xleslie3d_mpi+11595-111t loaded.'

When we bring up Reveal, we get the high level loops listed

CLICK ON IMPORTANT LOOP – RIGHT CLICK TO SCOPE

Navigation

Loop Performance

Time	Label
60.1097	LES3D@216
59.5277	LES3D@272
14.4611	FLUXI@21
14.4611	Instant
14.4605	FLUXI@22
12.7245	FLUXK@28
12.4870	FLUXJ@28
9.3702	FLUXJ@29
8.7140	FLUXK@29
6.2599	FLUXI@782
5.4454	FLUXJ@347
4.9909	FLUXK@348
4.9051	FLUXK@141
4.9044	FLUXK@143
4.7043	FLUXI@60
4.1214	FLUXJ@141
4.1210	FLUXJ@142
4.0092	FLUXK@69
3.9955	FLUXK@71
3.3823	FLUXI@123

Source - ... /lus/scratch/levesque/Leslie_CUG/leslie3d_mpi_reveal/src/fluxi.f

cce/8.7.10 Up Down Save

```
DO K = 1, KCMAX
DO J = 1, JCMAX
DO I = 0, ICMAX
  QS(I) = UAV(I,J,K) * SIX(I,J,K) +
  >      VAV(I,J,K) * SIY(I,J,K) +
  >      WAV(I,J,K) * SIZ(I,J,K)

  IF ( NSCHEME .EQ. 2 ) THEN
    L = I + 1 - IADD
    QSP = U(L,J,K) * SIX(I,J,K) +
    >      V(L,J,K) * SIY(I,J,K) +
    >      W(L,J,K) * SIZ(I,J,K)
    QSPI = (QSP - QS(I)) * DBLE(1 - 2 * IADD)
    IF ( QSPI .GT. 0.0D0 ) QS(I) = 0.5D0 * (QS(I) + QSP)
  ENDIF
  EST(I+1) = UAV(I+1,K) * QS(I)
```

Info - Line 21

- A loop starting at line 21 has been flattened (all calls inlined).
- A loop starting at line 21 was not vectorized because the target array (qs) would require random access.
- A loop starting at line 22 has been flattened (all calls inlined).

leslie3d.pl loaded. xleslie3d_mpi+11595-111t loaded.

Annotation of
optimization of code

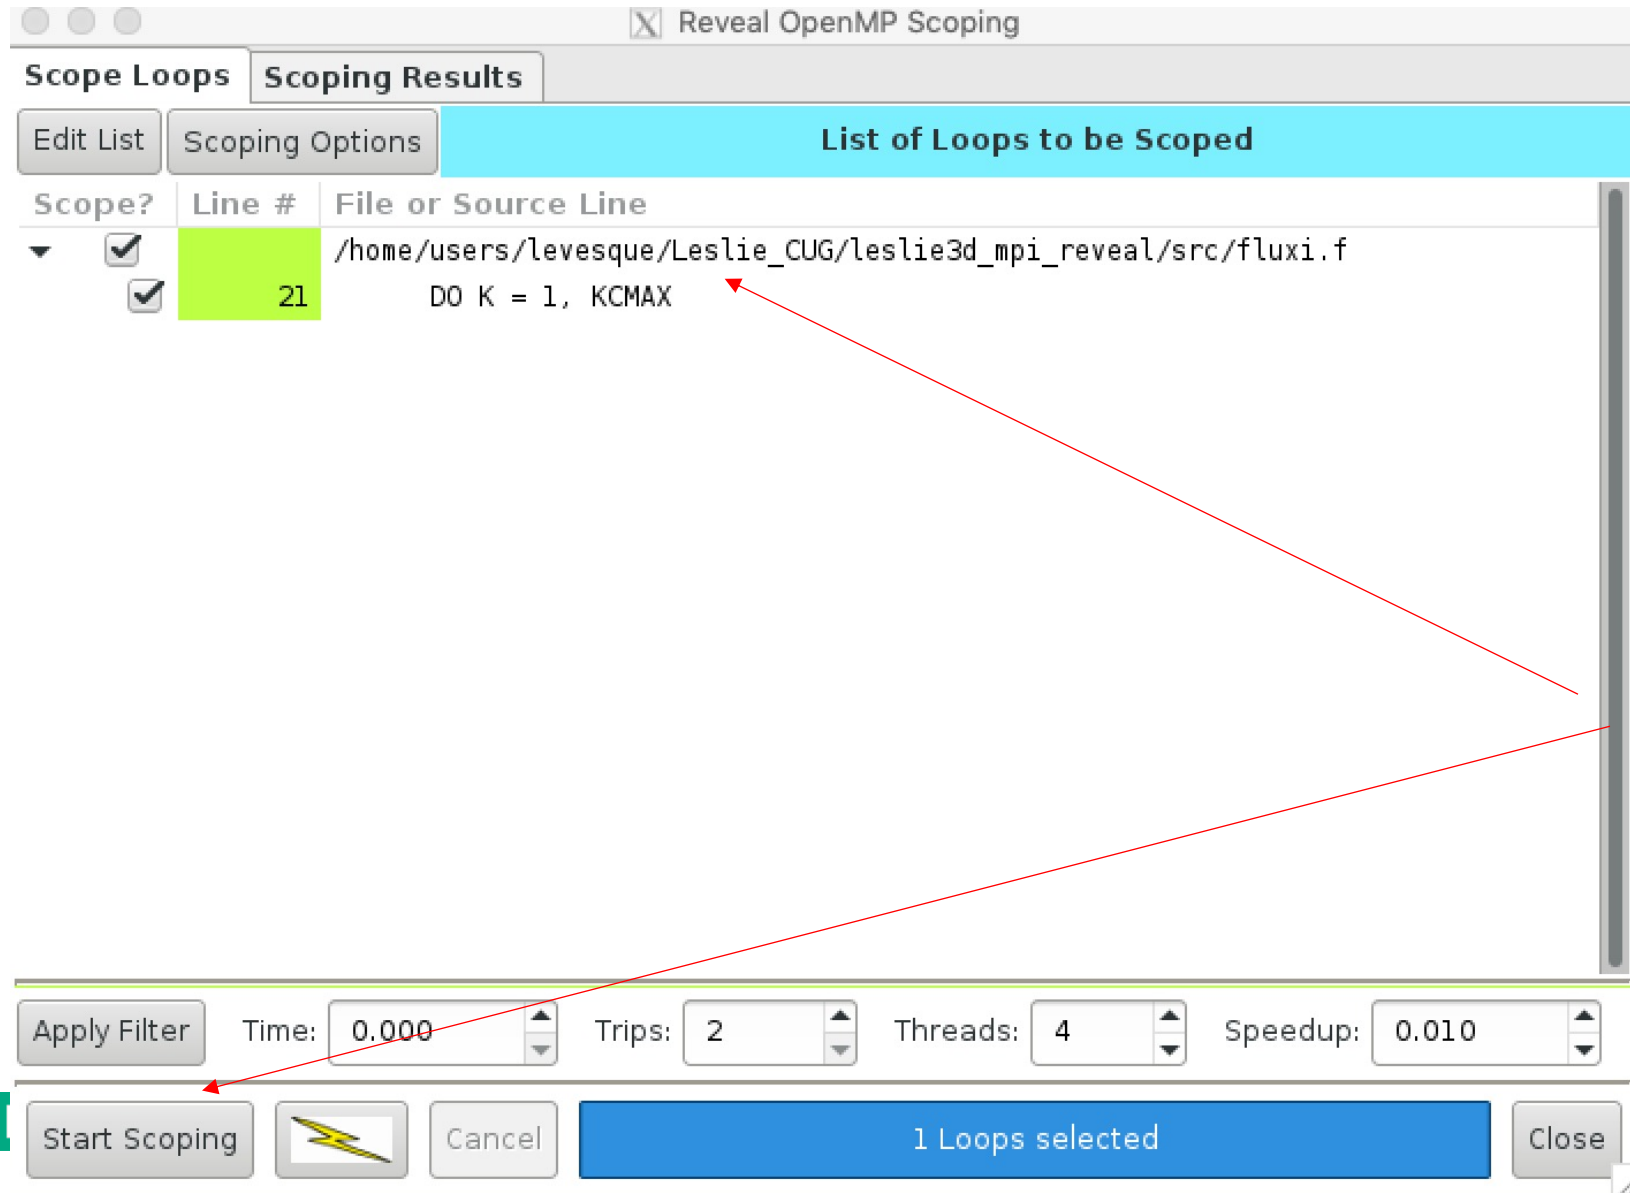
F- Flattened

V- Vectorized

Listing of loop

Diagnostics from
compiler

SCOPING WINDOW



Then up pops the
Scoping window

Loop selected Just click on
Start Scoping

SCOPING RESULTS

Name	Type	Scope	Info
fsi	Array	P S C <u>U</u>	WARN: LastPrivate of array may be very expensive. FAIL: FirstPrivate/Shared Scope Conflict.
i	Scalar	P S C <u>U</u>	FAIL: Possible recurrence involving this object.
icmax	Scalar	P S C <u>U</u>	FAIL: conflicting requirements, unable to scope.
l	Scalar	P S C <u>U</u>	FAIL: Possible recurrence involving this object.
qs	Array	P S C <u>U</u>	WARN: LastPrivate of array may be very expensive. FAIL: Last defining iteration not known for variable that may be live on exit.
j	Scalar	<u>P</u> S C U	
k	Scalar	<u>P</u> S C U	
qsp	Scalar	<u>P</u> S C U	
qspi	Scalar	<u>P</u> S C U	
dq	Array	P <u>S</u> C U	
dtv	Array	P <u>S</u> C U	
iadd	Scalar	P S C U	

First/Last Private: ☐ Enable FirstPrivate ☐ Enable LastPrivate

Reduction: None

Find Name:

P – Private

S – Shared

C – Conflict

U - Unresolved

User can now change scope by selecting the appropriate letter

ARRAY CONSTANTS ARE DIFFICULT TO SCOPE, ESPECIALLY WHEN PASSED TO A ROUTINE

The screenshot shows the Leslie3D MPI Reveal IDE. The main window displays the source code for `fluxi.f`. The variable `FLUXI@21` is highlighted in the navigation pane. The source code shows several occurrences of `FLUXI` (e.g., `FLUXI(I,5)`, `FLUXI(I,7)`, `FLUXI(I,L)`, `FLUXI@60`) and its use as an array constant in a function call `CALL VISCI (0, ICMAX, J, K, FLUXI)`. The IDE interface includes a navigation pane on the left, a main source code editor, and an information pane at the bottom.

Navigation

- Loop Performance
- 60.1097 LES3D@216
- 59.5277 LES3D@272
- 14.4611 FLUXI@21
- 14.4611 Instant
- 14.4605 FLUXI@22
- 12.7245 FLUXK@28
- 12.4870 FLUXJ@28
- 9.3702 FLUXJ@29
- 8.7140 FLUXK@29
- 6.2599 FLUXI@782
- 5.4454 FLUXJ@347
- 4.9909 FLUXK@348
- 4.9051 FLUXK@141
- 4.9044 FLUXK@143
- 4.7043 FLUXI@60
- 4.1214 FLUXJ@141
- 4.1210 FLUXJ@142
- 4.0092 FLUXK@69
- 3.9955 FLUXK@71
- 3.3823 FLUXI@123

Source - ... /lus/scratch/levesque/Leslie_CUG/leslie3d_mpi_reveal/src/fluxi.f

```
cce/8.7.10 Up Down Save
```

```
PAV(I,J,K) * SIZ(I,J,K)
    FLUXI(I,5) = (QAV(I,J,K,5) + PAV(I,J,K)) *
    QS(I)

    IF ( ISGSK .EQ. 1 ) THEN
        FLUXI(I,7) = QAV(I,J,K,7) * QS(I)
    ENDIF

    IF ( ICHEM .GT. 0 ) THEN
        DO L = 8, 7 + NSPECI
            FLUXI(I,L) = QAV(I,J,K,L) * QS(I)
        ENDDO
    ENDIF
ENDDO
    IF ( VISCOUS ) CALL VISCI ( 0, ICMAX, J, K, FLUXI )
```

Info - Line 21

- A loop starting at line 21 was scoped with errors. See Scoping Tool for more information.
- A loop starting at line 21 has been flattened (all calls inlined).
- A loop starting at line 21 was not vectorized because the target array (qs) would require ran

If you select a variable in the scoping window, all occurrences are highlighted in the main window

WHAT CAN THE USER DO TO HELP REVEAL

- Trace each variable that is unresolved and decide whether it is a potential race condition or can be scoped private or shared
 - If okay, make private or shared
- Once you have resolved all the unresolved variables then select the Insert Directives
- Couple definitions
 - Array Constant – An array not referenced by the parallel loop index
 - Array Constant reduction - example

With respect to K - $A(I,J)$ is an array reduction

```
do j=1,n
  do l=1,n
    do k=1,n
      a(i,j) = a(i,j) + b(i,k)*c(j,k)
    enddo
  enddo
enddo
```

SOME SIMPLE SCOPING RULES

- A scalar or an array not dependent on the loop being parallelized (array constant) should be private or ordered dependency
 - If the scalar is set prior to being used each time through the loop, then it is private. If it is not then it will result in a race condition
 - All elements of a array constant must be set prior to being used each pass through the loop. In those cases where not all the elements of the array are set prior to being used, first value getting is required.

SOME SIMPLE SCOPING RULES

- All arrays dependent upon the loop being parallelized should be shared. The compiler must perform data dependency analysis on the arrays to assure that there is no order dependency
- A reduction variable or constant array reduction is a special case that is identified by most compilers and must be in a reduction clause.
- Any variable that is just read is a shared variable

COMPLICATIONS DOWN THE CALL CHAIN AND MODULES

- No scoping directives can be inserted within those routines called from a parallelized loop.
 - All global variables must be shared
 - All variables allocated on stack must be private
- No private variables are allowed within a COMMON block or a module unless they are noted on a THREADPRIVATE directive

LASTPRIVATE SAVING AND FIRSTPRIVATE GETTING

- These two issues can drive you mad
 - Last Value saving is when the last value of a private variable is used outside the parallel loop. I can say that I have seen this in about .00001 % of the applications I have looked at – very rare.
 - When you select last value saving, the compiler has to make sure that the master thread either execute the last pass through the loop or have the thread, that does the last pass, copy its private variable to the master
 - Last value saving can have a big performance hit

THE COMPILER DOESN'T SCOPE EITHER FSI OR QS AS PRIVATE

The screenshot displays the xlc compiler interface with the file `leslie3d.pl` open. The **Navigation** pane on the left shows a list of code blocks with their performance metrics. The **Source** pane on the right shows the Fortran code for `fluxi.f`, with lines 21 through 36 highlighted in green, indicating they have been flattened. The **Info** pane at the bottom provides details about the scoping and vectorization of the code.

Navigation

Time	Label
60.1097	LES3D@216
59.5277	LES3D@272
14.4611	FLUXI@21
14.4611	Instant
14.4605	FLUXI@22
12.7245	FLUXK@28
12.4870	FLUXJ@28
9.3702	FLUXJ@29
8.7140	FLUXK@29
6.2599	FLUXI@782
5.4454	FLUXJ@347
4.9909	FLUXK@348
4.9051	FLUXK@141
4.9044	FLUXK@143
4.7043	FLUXI@60
4.1214	FLUXJ@141
4.1210	FLUXJ@142
4.0092	FLUXK@69
3.9955	FLUXK@71
3.3823	FLUXI@123

Source - ... /lus/scratch/levesque/Leslie_CUG/leslie3d_mpi_reveal/src/fluxi.f

```
22      DO J = 1, JCMAX
23      DO I = 0, ICMAX
24          QS(I) = UAV(I,J,K) * SIX(I,J,K) +
25              > VAV(I,J,K) * SIY(I,J,K) +
26              > WAV(I,J,K) * SIZ(I,J,K)
27
28      IF ( NSCHEME .EQ. 2 ) THEN
29          L = I + 1 - IADD
30          QSP = U(L,J,K) * SIX(I,J,K) +
31              > V(L,J,K) * SIY(I,J,K) +
32              > W(L,J,K) * SIZ(I,J,K)
33          QSPI = (QSP - QS(I)) * DBLE(1 - 2 * IADD)
34          IF ( QSPI .GT. 0.0D0 ) QS(I) = 0.5D0 * (QS(I) +
35              QSPI)
36      ENDIF
37      FSI(I,1) = QAV(I,J,K,1) * QS(I)
```

Info - Line 21—

- A loop starting at line 21 was scoped with errors. See Scoping Tool for more information.
- A loop starting at line 21 has been flattened (all calls inlined).
- A loop starting at line 21 was not vectorized because the target array (qs) would require random access.

leslie3d.pl loaded. xleslie3d_mpi+11595-1111t loaded.

AFTER ANALYZING THE VARIABLES THAT ARE UNRESOLVED

- If you are confident of your changes and all unresolved are resolved then click on INSERT Directives
 - CAUTION – if you are wrong you will be inserting a race condition
- If you cannot confidently resolve the unresolved, then go on to another loop



Navigation

Loop Performance

▶	333.8959	LES3D@216	★
▶	331.7217	LES3D@272	★
▶	89.0326	FLUXK@28	★
▶	82.6814	FLUXK@29	★
▶	81.3566	FLUXJ@28	★
▶	75.7702	FLUXJ@29	★
▶	75.4584	FLUXI@21	★
▶	75.4535	FLUXI@22	★
▶	63.5748	FLUXK@344	★
▶	57.5292	FLUXJ@340	★
▶	52.7949	FLUXI@782	★
▶	16.9245	FLUXI@128	★
▶	16.9218	FLUXI@129	★
▶	16.7026	FLUXK@138	★
▶	16.7000	FLUXK@140	★
▶	13.7771	FLUXJ@135	★
▶	13.7737	FLUXJ@136	★
▶	13.2113	UPDATE@10	★
▶	13.2083	UPDATE@11	★
▶	13.1524	PARALLEL@16	★
▶	12.7827	UPDATE@12	★
▶	6.7552	FLUXI@156	★
▶	6.7498	FLUXI@157	★
▶	6.6227	FLUXI@158	★
▶	6.3431	FLUXK@69	★
▶	6.0265	FLUXK@70	★
▶	5.5797	FLUXJ@69	★
▶	5.2100	FLUXI@62	★

Source - /lus/scratch/levesque/Video/leslie3d_mpi/src/fluxk.f

cce/8.7.5

Up

Down

Save

```
+
!1 26 CALL EXTRAPK ( FSK, I )
27
! Directive inserted by Cray Reveal. May be incomplete.
!$OMP parallel do default(none)
!$OMP& private (fsk,i,j,k,l,qs,qsp,qspk)
!$OMP& shared (dq,dtv,ind,jcmax,kadd,kcmax,pav,qav,skx,sky,skz,u,uav,
!$OMP& v,vav,w,wav)
FS 28 DO J = 1, JCMAX
F 29 DO K = 0, KCMAX
30
FV 31 QS(1:IND) = UAV(1:IND,J,K) * SKX(1:IND,J,K) +
32 > VAV(1:IND,J,K) * SKY(1:IND,J,K) +
33 > WAV(1:IND,J,K) * SKZ(1:IND,J,K)
34
35 IF ( NScheme .EQ. 2 ) THEN
36 L = K + 1 - KADD
37 DO I = 1, IND
38 QSP = U(I,J,L) * SKX(I,J,K) +
39 > V(I,J,L) * SKY(I,J,K) +
40 > W(I,J,L) * SKZ(I,J,K)
41 QSPK = (QSP - QS(I)) * DBLE(1 - 2 * KADD)
42 IF (QSPK .GT. 0.0D+00) QS(I) = 0.5D+00 * (QS(I) + QSP)
43 ENDDO
44 ENDIF
45
```

Info

REVEAL HELPED ANALYZE THIS LOOP

```
28.                ! Directive inserted by Cray Reveal.  May be incomplete.
29.      M-----< !$OMP  parallel do default(none)
30.      M                !$OMP&  private (fsk,i,j,k,l,qs,qsp,qspk)
31.      M                !$OMP&  shared  (dq,dtv,ind,jcmax,kadd,kcmax,pav,qav,skx,sky,skz,u,uav,
32.      M                !$OMP&                                v,vav,w,wav)
33. + M mF-----<      DO J = 1, JCMAX
34. + M mF F-----<      DO K = 0, KCMAX
35.      M mF F V-----<      DO I = 1, IND
36.      M mF F V                QS(I) = UAV(I,J,K) * SKX(I,J,K) +
37.      M mF F V                >          VAV(I,J,K) * SKY(I,J,K) +
38.      M mF F V                >          WAV(I,J,K) * SKZ(I,J,K)
39.      M mF F V
40.      M mF F V                IF ( NScheme .EQ. 2 ) THEN
41.      M mF F V                L = K + 1 - KADD
42.      M mF F V
43.      M mF F V
44.      M mF F V
45.      M mF F V
46.      M mF F V                IF (QSPK .GT. 0.0D+00) QS(I) = 0.5D+00 * (QS(I) + QSP)
47.      M mF F V                ENDIF
48.      M mF F V
49.      M mF F V                FSK(I,K,1) =  QAV(I,J,K,1) * QS(I)
50.      M mF F V                FSK(I,K,2) =  QAV(I,J,K,2) * QS(I) +
51.      M mF F V                >          PAV(I,J,K) * SKX(I,J,K)
52.      M mF F V                FSK(I,K,3) =  QAV(I,J,K,3) * QS(I) +
53.      M mF F V                >          PAV(I,J,K) * SKY(I,J,K)
54.      M mF F V                FSK(I,K,4) =  QAV(I,J,K,4) * QS(I) +
55.      M mF F V                >          PAV(I,J,K) * SKZ(I,J,K)
56.      M mF F V                FSK(I,K,5) = (QAV(I,J,K,5) + PAV(I,J,K)) *
57.      M mF F V                >          QS(I)
```

Notice that Reveal doesn't use default shared, primarily to illustrate that it has handled all variables

OPENMP VERSUS MPI

Running on 7 nodes with 48 threads

Table 1: Profile by Function

Samp%	Samp	Imb. Samp	Imb. Samp%	Group Function=[MAX10] PE=HIDE Thread=HIDE
100.0%	2,038.3	--	--	Total
48.6%	990.9	--	--	USER
11.3%	231.0	8.0	3.9%	tmstep_
6.0%	122.0	10.0	8.8%	parallel_
5.7%	115.3	28.7	23.3%	ghost_
4.0%	81.9	4.1	5.6%	fluxi_.LOOP@li.30
3.6%	72.7	6.3	9.3%	fluxk_.LOOP@li.38
3.6%	72.6	7.4	10.8%	fluxj_.LOOP@li.38
3.5%	71.9	85.1	63.3%	wallbc_
42.3%	862.1	--	--	MPI
17.3%	352.9	27.1	8.3%	MPI_REDUCE
16.3%	332.1	69.9	20.3%	MPI_SEND
5.9%	121.1	30.9	23.7%	MPI_ALLREDUCE
7.5%	152.3	--	--	ETC
1.6%	32.6	--	--	OMP

Running on 7 nodes with 1 threads

Table 1: Profile by Function

Samp%	Samp	Imb. Samp	Imb. Samp%	Group Function=[MAX10] PE=HIDE
100.0%	9,623.4	--	--	Total
79.7%	7,667.6	--	--	USER
20.5%	1,971.0	302.0	15.5%	fluxi_.LOOP@li.30
18.8%	1,812.7	427.3	22.3%	fluxk_.LOOP@li.38
17.5%	1,688.4	230.6	14.0%	fluxj_.LOOP@li.38
5.2%	496.1	30.9	6.8%	extrapk_.LOOP@li.158
4.6%	438.0	33.0	8.2%	extrapj_.LOOP@li.159
4.3%	414.4	5.6	1.5%	update_.LOOP@li.18
2.4%	229.0	6.0	3.0%	tmstep_
19.5%	1,878.0	--	--	MPI
8.7%	835.7	789.3	56.7%	MPI_SEND
5.7%	550.4	802.6	69.2%	MPI_WAIT
3.7%	354.4	26.6	8.1%	MPI_REDUCE

OPENMP VERSUS MPI

Running on 7 nodes with 48 threads

Table 1: Profile by Function

Samp%	Samp	Imb. Samp	Imb. Samp%	Group Function=[MAX10] PE=HIDE Thread=HIDE
100.0%	2,038.3	--	--	Total
48.6%	990.9	--	--	USER
11.3%	231.0	8.0	3.9%	tmstep_
6.0%	122.0	10.0	8.8%	parallel_
5.7%	115.3	28.7	23.3%	ghost_
4.0%	81.9	4.1	5.6%	fluxi_.LOOP@li.30
3.6%	72.7	6.3	9.3%	fluxk_.LOOP@li.38
3.6%	72.6	7.4	10.8%	fluxj_.LOOP@li.38
3.5%	71.9	85.1	63.3%	wallbc_
42.3%	862.1	--	--	MPI
17.3%	352.9	27.1	8.3%	MPI_REDUCE
16.3%	332.1	69.9	20.3%	MPI_SEND
5.9%	121.1	30.9	23.7%	MPI_ALLREDUCE
7.5%	152.3	--	--	ETC
1.6%	32.6	--	--	OMP

Running on 7 nodes with 224 MPI tasks*

Table 1: Profile by Function

Samp%	Samp	Imb. Samp	Imb. Samp%	Group Function=[MAX10] PE=HIDE
100.0%	6,633.3	--	--	Total
92.2%	6,114.0	--	--	MPI
48.9%	3,246.3	1,556.7	32.6%	MPI_ALLREDUCE
22.9%	1,521.1	3,025.9	66.8%	MPI_REDUCE
11.8%	785.4	530.6	40.5%	MPI_SEND
7.8%	518.0	579.0	53.0%	MPI_WAIT
6.9%	459.7	--	--	USER
1.7%	112.1	70.9	38.9%	fluxi_.LOOP@li.30
1.3%	88.4	25.6	22.6%	fluxk_.LOOP@li.38
1.2%	81.2	33.8	29.5%	fluxj_.LOOP@li.38

* When I ran 336 Mpi tasks – it seg faulted

WHY IS MPI GETTING BETTER SCALING THAN OPENMP ON THE NODE

- OpenMP run
 - `setenv OMP_NUM_THREADS 48`
 - `srun -n 7 -N7 ./xleslie3d_mpi`
- MPI run
 - `setenv OMP_NUM_THREADS 1`
 - `srun -n 336 -N7 ./xleslie3d_mpi`
- Improved OpenMP
 - `setenv OMP_NUM_THREADS 24`
 - `srun -n 14 -N7 ./xleslie3d_mpi`
 - `setenv OMP_NUM_THREADS 12`
 - `srun -n 28 -N7 ./xleslie3d_mpi`

Running across four NUMA domains will incur NUMA issues

Running with a MPI task on each NUMA domain and 12 threads/MPI task

SRUN -N 7 -N7 OMP_NUM_THREADS = 48

Table 4: Memory Bandwidth by Numanode

Memory Traffic GBytes	Read Memory Traffic GBytes	Write Memory Traffic GBytes	Thread Time	Memory Traffic GBytes / Sec	Memory Traffic / Nominal Peak	Numanode Node Id=[max3,min3] PE=HIDE Thread=HIDE
692.02	546.44	145.58	20.905607	33.10	12.9%	numanode.0
692.02	546.44	145.58	20.826853	33.23	13.0%	nid.19
660.64	523.35	137.29	20.905607	31.60	12.3%	nid.23
636.23	502.95	133.28	20.851091	30.51	11.9%	nid.17
98.04	79.26	18.78	14.413192	6.80	2.7%	nid.21
52.58	42.79	9.80	14.388447	3.65	1.4%	nid.22
537.23	424.03	113.20	20.827913	25.79	10.1%	numanode.1
537.23	424.03	113.20	20.827913	25.79	10.1%	nid.18
289.46	229.36	60.10	14.449583	20.03	7.8%	nid.20
49.34	40.25	9.09	14.358161	3.44	1.3%	nid.19
34.87	28.08	6.79	14.431710	2.42	0.9%	nid.23
492.50	388.74	103.76	20.828108	23.65	9.2%	numanode.2
492.50	388.74	103.76	20.826837	23.65	9.2%	nid.22
416.20	328.86	87.34	20.828108	19.98	7.8%	nid.20
414.88	327.01	87.87	20.817194	19.93	7.8%	nid.21
80.90	64.72	16.18	14.427566	5.61	2.2%	nid.17
41.14	33.01	8.13	14.378477	2.86	1.1%	nid.18
207.95	164.01	43.94	14.437465	14.40	5.6%	numanode.3

SRUN -N 14 -N7 OMP_NUM_THREADS = 24

able 4: Memory Bandwidth by Numanode

Memory Traffic GBytes	Read Memory Traffic GBytes	Write Memory Traffic GBytes	Thread Time	Memory Traffic GBytes / Sec	Memory Traffic / Nominal Peak	Numanode Node Id=[max3,min3] PE=HIDE Thread=HIDE
374.82	302.81	72.01	17.936564	20.90	8.2%	numanode.0
374.82	302.81	72.01	17.918537	20.92	8.2%	nid.17
368.30	297.39	70.91	17.919944	20.55	8.0%	nid.22
364.44	294.10	70.33	17.936564	20.32	7.9%	nid.20
50.46	38.63	11.83	17.918304	2.82	1.1%	nid.19
29.65	25.08	4.57	11.629136	2.55	1.0%	nid.23
4.53	3.78	0.75	11.629711	0.39	0.2%	nid.18
356.36	287.94	68.42	17.943760	19.86	7.8%	numanode.1
356.36	287.94	68.42	17.916977	19.89	7.8%	nid.19
348.44	281.47	66.97	17.942056	19.42	7.6%	nid.18
343.78	277.59	66.19	17.933473	19.17	7.5%	nid.17
318.21	255.83	62.38	17.943760	17.73	6.9%	nid.23
4.24	3.56	0.68	11.611935	0.37	0.1%	nid.20
376.72	303.36	73.37	17.940639	21.00	8.2%	numanode.2
376.72	303.36	73.37	17.935383	21.00	8.2%	nid.23
371.79	300.44	71.35	17.926640	20.74	8.1%	nid.21
367.23	295.95	71.28	17.921852	20.49	8.0%	nid.22
355.70	288.58	67.13	11.620607	30.61	12.0%	nid.19
351.35	282.50	68.84	17.938276	19.59	7.7%	nid.20
34.92	29.00	5.91	11.593323	3.01	1.2%	nid.17
367.95	296.27	71.69	17.937994	20.51	8.0%	numanode.3

SRUN -N 28 -N7 OMP_NUM_THREADS = 12

Table 4: Memory Bandwidth by Numanode

Memory Traffic GBytes	Read Memory Traffic GBytes	Write Memory Traffic GBytes	Thread Time	Memory Traffic GBytes / Sec	Memory Traffic / Nominal Peak	Numanode Node Id=[max3,min3] PE=HIDE Thread=HIDE
228.24	183.59	44.65	16.163097	14.12	5.5%	numanode.0
228.24	183.59	44.65	16.148021	14.13	5.5%	nid.17
223.37	180.53	42.84	16.163097	13.82	5.4%	nid.22
198.67	160.15	38.51	16.155601	12.30	4.8%	nid.23
193.59	156.32	37.27	16.159357	11.98	4.7%	nid.18
191.25	154.41	36.83	16.159708	11.83	4.6%	nid.20
190.05	154.30	35.75	16.162672	11.76	4.6%	nid.21
234.78	187.65	47.14	16.175862	14.51	5.7%	numanode.1
233.23	187.65	45.58	16.166679	14.43	5.6%	nid.18
232.62	185.48	47.14	16.175862	14.38	5.6%	nid.22
205.87	165.90	39.97	16.161616	12.74	5.0%	nid.19
200.60	161.63	38.97	16.154478	12.42	4.9%	nid.17
156.37	127.68	28.69	9.974652	15.68	6.1%	nid.20
156.16	127.78	28.38	9.981468	15.64	6.1%	nid.21
234.62	187.48	47.13	16.176280	14.50	5.7%	numanode.2
234.62	187.48	47.13	16.167319	14.51	5.7%	nid.21
232.10	185.04	47.05	16.176280	14.35	5.6%	nid.20
224.69	181.56	43.12	16.176174	13.89	5.4%	nid.18
200.39	162.84	37.55	16.152474	12.41	4.8%	nid.17
198.10	160.08	38.03	16.160609	12.26	4.8%	nid.19
156.38	127.36	29.02	9.975610	15.68	6.1%	nid.22
198.06	160.32	37.74	16.171106	12.25	4.8%	numanode.3

SRUN -N 336 -N7 OMP_NUM_THREADS = 1

Table 3: Memory Bandwidth by Numanode

Memory Traffic GBytes	Read Memory Traffic GBytes	Write Memory Traffic GBytes	Thread Time	Memory Traffic GBytes / Sec	Memory Traffic / Nominal Peak	Numanode Node Id=[max3,min3] PE=HIDE
339.06	282.68	56.38	72.507205	4.68	1.8%	numanode.0
339.06	282.68	56.38	71.368068	4.75	1.9%	nid.18
332.41	276.58	55.83	70.427465	4.72	1.8%	nid.21
298.48	250.14	48.35	70.281285	4.25	1.7%	nid.22
284.96	239.95	45.02	72.507205	3.93	1.5%	nid.17
257.74	215.10	42.63	70.606828	3.65	1.4%	nid.20
253.79	211.00	42.78	69.717114	3.64	1.4%	nid.23
298.16	246.34	51.82	73.258873	4.07	1.6%	numanode.1
298.16	246.34	51.82	73.258873	4.07	1.6%	nid.17
291.14	244.89	46.25	70.993035	4.10	1.6%	nid.20
264.68	220.22	44.47	69.997530	3.78	1.5%	nid.22
255.54	211.79	43.75	69.723811	3.67	1.4%	nid.23
255.31	213.63	41.68	70.415287	3.63	1.4%	nid.21
222.18	184.48	37.71	70.866074	3.14	1.2%	nid.19
317.07	267.29	49.78	72.149334	4.39	1.7%	numanode.2
317.07	267.29	49.78	69.717081	4.55	1.8%	nid.23
301.05	251.67	49.38	70.680637	4.26	1.7%	nid.19
299.67	250.23	49.44	70.813903	4.23	1.7%	nid.20
294.69	246.52	48.17	71.229050	4.14	1.6%	nid.18
265.58	220.85	44.74	70.201836	3.78	1.5%	nid.21
248.40	208.18	40.22	72.149334	3.44	1.3%	nid.17
367.85	307.58	60.27	73.012099	5.04	2.0%	numanode.3